

Introduction to Python

Christopher Barker

UW Continuing Education / Isilon

July 25, 2012

Table of Contents

- 1 Review/Questions
- 2 Sockets
- 3 HTTP
- 4 Building an HTTP server
- 5 Wrap Up / Homework

Review of Previous Class

- Keyword arguments/parameters
- Lists
- Dictionaries
- Sets

Homework review

Dave Thomas Coding Kata 14

http:

[//codekata.pragprog.com/2007/01/kata_fourteen_t.html](http://codekata.pragprog.com/2007/01/kata_fourteen_t.html)

My results:

honour my gems hundred year old forever do not even two might
end was a he realised how yes when i stronger if i perfectly happy
and revealed it to windows were blocked a thing is to sound him rd
floor that would induce the work by people must sit without with
jabez wilson bad compliment ...

My code: `week-04/code/trigram.py`

Lightning Talk

Lightning Talk:

David

Class Structure

This class is different – more a tutorial than a class:
lots of coding.

We're going to run through building a really basic
HTTP server from the ground up.

We'll see how far we get.

Note: I'm no expert – I'm learning along with you...

Sockets

“Socket” at either end of a pathway: client and server can be “plugged in” to communicate

Five pieces of data to uniquely identify a connection:

- Transport protocol (UDP, TCP) (we'll use TCP)
- Remote IP address
- Remote port number
- Local IP address
- Local port number

(use localhost (127.0.0.1) on both ends for this class...)

Python Socket Module

Create a socket:

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

AF_INET : Internet Family of Protocols

SOCK_STREAM : TCP

Set an option:

```
s.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
```

re-use the address – so the OS won't reserve it

(Python docs say "see the UNIX man pages...")

A socket server

```
echo_server.py
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
```

```
# the bind makes it a server
```

```
s.bind( (host, port) )
```

```
s.listen(backlog)
```

```
while True: # keep looking for new connections forever
```

```
    client, address = s.accept() # look for a connection
```

```
    data = client.recv(size)
```

```
    if data: # if the connection was closed there would be
```

```
        print "received: %s, sending it back"%data
```

```
        client.send(data)
```

```
        client.close()
```

A socket client

echo_client.py

```
while True:
    s = socket.socket(socket.AF_INET,
                      socket.SOCK_STREAM)
    s.connect((host,port))
    msg = raw_input('what should I send? >> ')
    if msg:
        s.send(msg)
        data = s.recv(size)
        s.close()
```

Mini-LAB

Start up `echo_server.py`

Start up `echo_client.py`

(in different terminals...)

Watch what happens when you use the client

If any of you are using a shared system – change your port numbers

NOTE: running from iPython can cause trouble...

Lightning Talk

Lightning Talk:

Drew

HTTP

HyperText Transfer Protocol Client-Server:

- requests
- responses

Each has:

- Method specification (request)
- Status line (response)
- Headers (RFC 822-compliant)
(optionally)
- Entity headers and body

(RFC 2616)

HTTP Requests

Request Methods

- GET – Read URI content
- HEAD – GET headers only
- POST – Create
- PUT – Update (entity transfer to server)
- DELETE – delete content

(GET and POST different ways to do similar things...)

There are four others – but these are the ones most used

HTTP request

Example HTTP GET request

GET /a_file HTTP/1.1

Host: localhost:55555

User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.7; rv

Accept: text/html,application/xhtml+xml,application/xml;q=0

Accept-Language: en-us,en;q=0.5

Accept-Encoding: gzip, deflate

Connection: keep-alive

HTTP Responses

Response Codes

- 200 OK
- 404 Not Found
- 301 Moved Permanently
- 302 Moved Temporarily
- 303 See Other (HTTP 1.1 only)
- 500 Server Error

There are four others – but these are the ones most used

HTTP Response header

```
HTTP/1.1 200 OK
```

```
Date: Fri, 31 Dec 1999 23:59:59 GMT
```

```
Content-Type: text/html
```

```
Content-Length: 1354
```

```
<html>
```

```
<body>
```

```
<h1>Happy New Millennium!</h1>
```

```
(more file contents)
```

```
... </body> </html>
```

Blank line between header and body critical!

(\r\n linefeeds)

HTTP Response header

Header-Name: value

Quick reference to HTTP headers:

`http://www.cs.tut.fi/~jkorpela/http.html`

HTTP Response header

body data:

Content-Type: xyz/abc

Mime types we might want:

- text/plain
- text/html
- image/png
- image/jpeg

<http://www.webmaster-toolkit.com/mime-types.shtml>

Debugging

Debugging Tools

- windows:
`http://www.fiddler2.com/fiddler2/`
- windows & mac:
`http://www.charlesproxy.com/`
- Firefox:
`http://getfirebug.com/`
- Safari, Chrome and IE: built in

Building an HTTP Server

We've got everything we need to know to build a simple server

(GET only for now...)

Build an HTTP server that can serve up the files in: `week-05\code\web`

Building an HTTP Server

Incremental Development:

- 1 A socket server that can receive a request (and print that request to the console)
- 2 Server returns a simple reply
- 3 Server returns a properly formatted HTTP reply
- 4 Server returns a 404 error
- 5 Server returns the file asked for
- 6 Server returns a directory listing
- 7 Server returns multiple file types
- 8 Server returns a calculated response

http_serve1.py

Edit `echo_serve1.py` to print the request:

- 1 Point your browser at `echo_server.py` – what do you get?
- 2 Save it as `http_serve1.py`
- 3 Edit it to print the request to the console
- 4 Edit it to return a bit of html (`tiny_html.html`)
- 5 What happens when you point your browser to it?
Try a couple different browsers – I get a different result with Firefox and Safari

http_serve2.py

Return a properly formatted HTTP response:

- 1 Save `http_serve1.py` as `http_serve2.py`
- 2 Add code that generates an HTTP “200 OK” header (don’t forget the blank line! (`\r\n`))
- 3 Use `httpdate.httpdate_now()` to give you an HTTP date string
- 4 What happens when you point your browser to it now?

Lightning Talk

Lightning Talk:

Phillip

http_serve3.py

Parse the request:

- 1 Save `http_serve2.py` as `http_serve3.py`
- 2 Add code that parses the HTTP request – it should give you the URI requested
- 3 Have it check to make sure it's a GET request
- 4 print the URI (file name) to the console

http_serve4.py

Return a Listing:

- 1 Save `http_serve3.py` as `http_serve4.py`
- 2 Add code that parses the URI – so you can figure out what file is requested
- 3 check to see if is a directory or a file
- 4 return a listing (simple text) of the dir if it's a dir
- 5 return a 404 otherwise

http_serve5.py

Support various file types:

- 1 Save `http_serve4.py` as `http_serve5.py`
- 2 If the request is for a file – return that file
- 3 have it be a different mime type depending on the type of file
- 4 support: `.html`, `.txt`, `.jpeg`, `.png`

You now have a pretty functional web server!

http_serve6.py

If we have time...

- 1 Save `http_serve5.py` as `http_serve6.py`
- 2 Format the Dir listing as HTML
- 3 Make the files in the listing clickable links.

and / or

- 1 Make a simple web app (non-static)
- 2 Have `localhost:50000/the_time` return a web page with the current time.

You now have a very functional web server!

Standard Library Support

It's unlikely that you'll need to use raw sockets

Standard Library Modules

- `httplib`
- `urllib2` (requests (PyPI))
- `smtplib`
- `poplib`
- `imaplib`

Third Party Modules

- `requests` (cleaner interface than `urllib2`)
- `paramiko` (SSH)
- Probably one for anything you're likely to do...

Homework

- Think Python ch. 15-18
- Finish what you didn't get to in class
- When a *.py file is asked for:
 - Don't return the script's contents
 - Return the result of running the script (stdout)
(`subprocess.Popen()` will be useful)
 - Test with a script that prints the time (or something...)
 - You've just re-invented CGI